

Lab 3

Import necessary libraries

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from sklearn.datasets import load_diabetes
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, roc_curve, auc
```

Load the diabetes dataset

```
diabetes = load_diabetes()
```

```
X, y = diabetes.data, diabetes.target
```

Convert the target variable to binary (1 for diabetes, 0 for no diabetes)

```
y_binary = (y > np.median(y)).astype(int)
```

```
# Split the data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(
```

```
    X, y_binary, test_size=0.2, random_state=42)
```

Standardize features

```
scaler = StandardScaler()
```

```
X_train = scaler.fit_transform(X_train)
```

```
X_test = scaler.transform(X_test)
```

Train the Logistic Regression model

```
model = LogisticRegression()
```

```
model.fit(X_train, y_train)
```

Evaluate the model

```
y_pred = model.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print("Accuracy: {:.2f}%".format(accuracy * 100))
```

output:

Accuracy: 73.03%

evaluate the model

```
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

```
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

output:

Confusion Matrix:

```
[[36 13]
```

```
[11 29]]
```

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.77	0.73	0.75	49
---	------	------	------	----


1	0.69	0.72	0.71	40
---	------	------	------	----

```
accuracy          0.73    89
macro avg    0.73    0.73    0.73    89
weighted avg    0.73    0.73    0.73    89
```

Visualize the decision boundary with accuracy information

```
plt.figure(figsize=(8, 6))
sns.scatterplot(x=X_test[:, 2], y=X_test[:, 8], hue=y_test, palette={
    0: 'blue', 1: 'red'}, marker='o')
plt.xlabel("BMI")
plt.ylabel("Age")
plt.title("Logistic Regression Decision Boundary\nAccuracy: {:.2f}%".format(
    accuracy * 100))
plt.legend(title="Diabetes", loc="upper right")
plt.show()
```

output:



Plot ROC Curve

```
y_prob = model.predict_proba(X_test)[: , 1]
```

```
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
```

```
roc_auc = auc(fpr, tpr)
```

```
plt.figure(figsize=(8, 6))
```

```
plt.plot(fpr, tpr, color='darkorange', lw=2,
```

```
label=f'ROC Curve (AUC = {roc_auc:.2f})')
```

```
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--', label='Random')
```

```
plt.xlabel('False Positive Rate')
```

```
plt.ylabel('True Positive Rate')

plt.title('Receiver Operating Characteristic (ROC) Curve\nAccuracy: {:.2f}%'.format(
    accuracy * 100))

plt.legend(loc="lower right")

plt.show()
```

output:

